

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
```

```
In [2]: def print_title(title):
        print(f'\n{"-"*60}\n\033[1m{title}\033[0m')
def print_section(title):
    print(f'{"-"*60}\n{title}\n{"-"*60}')
```

1. Loading and Preprocessing

```
In [4]: # importing and reaing the data
cp = pd.read_csv('CarPrice.csv')
cp=pd.DataFrame(cp)
cp.head(5)
```

Out[4]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wh
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	

5 rows × 26 columns

```
In [5]: #Data frme Description
cp.describe()
```

Out[5]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	borerati
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317	3.32975
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693	0.27084
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000	2.54000
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000	3.15000
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000	3.31000
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000	3.58000
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000	3.94000

```
In [6]: # DataFrame info
cp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration             205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel            205 non-null   object
8   enginelocation         205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm               205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg            205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [7]: # Null value finding
cp.isnull().sum()
```


```
Out[7]: car_ID                0
symboling              0
CarName               0
fueltype              0
aspiration             0
doornumber            0
carbody               0
drivewheel            0
enginelocation         0
wheelbase             0
carlength             0
carwidth              0
carheight             0
curbweight            0
enginetype            0
cylindernumber        0
enginesize            0
fuelsystem            0
boreratio             0
stroke                0
compressionratio      0
horsepower            0
peakrpm               0
citympg               0
highwaympg            0
price                 0
dtype: int64
```

```
In [8]: # Datafarama Duplicate row finding
```

```
cp.duplicated().sum()
```

Out[8]: 0

```
In [9]: # Visualizing outliers of dataset before processing
plt.figure(figsize=(6,5))
cp.boxplot()
plt.xticks(rotation=90)
plt.show()
```



```
In [10]: ## finding object-type Features columns in dataframe and this will be compayer after enc
object_cols = [col for col in cp.columns if cp[col].dtype == 'object']
object_cols_df = pd.DataFrame(object_cols)
object_cols_df.columns = ['Object_column']
print(f'Object data type columns:\n{object_cols_df}\n\nThe total count of Object Columns
```

Object data type columns:

```
Object_column
0      CarName
1     fueltype
2    aspiration
3   doornumber
4     carbody
5   drivewheel
6  enginelocation
7    enginetype
8  cylindernumber
9    fuelsystem
```

The total count of Object Columns is 10

```
In [11]: print(f'total num of unique value in CarName: {len(cp['CarName'].unique())}')
print(f'total num of unique value in fueltype: {len(cp['fueltype'].unique())}, {(cp['fueltype'].unique())}')
print(f'total num of unique value in aspiration: {len(cp['aspiration'].unique())} {(cp['aspiration'].unique())}')
print(f'total num of unique value in doornumber: {len(cp['doornumber'].unique())} {(cp['doornumber'].unique())}')
print(f'total num of unique value in carbody: {len(cp['carbody'].unique())} {(cp['carbody'].unique())}')
print(f'total num of unique value in drivewheel: {len(cp['drivewheel'].unique())} {(cp['drivewheel'].unique())}')
print(f'total num of unique value in enginelocation: {len(cp['enginelocation'].unique())} {(cp['enginelocation'].unique())}')
print(f'total num of unique value in enginetype: {len(cp['enginetype'].unique())} {(cp['enginetype'].unique())}')
print(f'total num of unique value in cylindernumber: {len(cp['cylindernumber'].unique())} {(cp['cylindernumber'].unique())}')
print(f'total num of unique value in fuelsystem: {len(cp['fuelsystem'].unique())} {(cp['fuelsystem'].unique())}')

total num of unique value in CarName: 147
total num of unique value in fueltype: 2, ['gas' 'diesel']
total num of unique value in aspiration: 2 ['std' 'turbo']
total num of unique value in doornumber: 2 ['two' 'four']
total num of unique value in carbody: 5 ['convertible' 'hatchback' 'sedan' 'wagon' 'hardtop']
total num of unique value in drivewheel: 3 ['rwd' 'fwd' '4wd']
total num of unique value in enginelocation: 2 ['front' 'rear']
total num of unique value in enginetype: 7 ['dohc' 'ohcv' 'ohc' 'l' 'rotor' 'ohcf' 'dohcv']
total num of unique value in cylindernumber: 7 ['four' 'six' 'five' 'three' 'twelve' 'two' 'eight']
total num of unique value in fuelsystem: 8 ['mpfi' '2bbl' 'mfi' '1bbl' 'spfi' '4bbl' 'idi' 'spdi']
```

OrdinalEncoding

OrdinalEncoding method used for having more control over mapping

```
In [13]: # Encoding for 'cylindernumber' and 'doornumber' Feature value
from category_encoders import OrdinalEncoder
Enco_cp = cp
mapping = [{'col': 'cylindernumber', 'mapping': {'two': 2, 'four': 4, 'six': 6, 'five': 5, 'three': 3, 'eight': 8}}]
encoder = OrdinalEncoder(mapping=mapping)
Enco_cp = encoder.fit_transform(Enco_cp)
```

```
In [14]: # create dictionary for OrdinalEncoder
uniqu = Enco_cp['CarName'].unique()
print(f'The cont of unique name in uniqu df is {len(uniqu)} and count of unique name in original df is {len(cp['CarName'].unique())}')
print()
uniqu_df = pd.DataFrame(uniqu)
uniqu_df = uniqu_df.reset_index(drop=True)
uniqu_df.columns = ['Uniqu_CarName']
uniqu_df['CarName_Enco'] = uniqu_df.index + 1
# Create the dictionary
map_dict = pd.Series(uniqu_df['CarName_Enco'].values, index=uniqu_df['Uniqu_CarName']).to_dict()

The cont of unique name in uniqu df is 147 and count of unique name in original df of encoding is 147
```

```
In [15]: # Encoding for 'CarName'
mapping = [{'col': 'CarName', 'mapping': map_dict}]
encoder = OrdinalEncoder(mapping=mapping)
Enco_cp = encoder.fit_transform(Enco_cp)
Enco_cp = Enco_cp.rename(columns={'CarName': 'Car_Name'}) # renamed to avoid repeated columns
```

Custom mapping is done for all remaining object-type columns

```
In [17]: # custom mapping for fueltype
name_mapping = {'gas': 1, 'diesel': 2}
Enco_cp['fueltype'] = Enco_cp['fueltype'].map(name_mapping)
```

```

In [18]: # custom mapping for aspiration
name_mapping1 = {'std': 1, 'turbo': 2}
Enco_cp['aspiration'] = Enco_cp['aspiration'].map(name_mapping1)

In [19]: # custom mapping for carbody
name_mapping2 = {'convertible':1, 'hatchback':2, 'sedan':3, 'wagon':3, 'hardtop':4}
Enco_cp['carbody'] = Enco_cp['carbody'].map(name_mapping2)

In [20]: # custom mapping for drivewheel
name_mapping3 = {'rwd':1, 'fwd':2, '4wd':3}
Enco_cp['drivewheel'] = Enco_cp['drivewheel'].map(name_mapping3)

In [21]: # custom mapping for enginelocation
name_mapping4 = {'front':1, 'rear':2}
Enco_cp['enginelocation'] = Enco_cp['enginelocation'].map(name_mapping4)

In [22]: # custom mapping for enginetype
name_mapping5 = {'dohc':1, 'ohcv':2, 'ohc':3, 'l':3, 'rotor':4, 'ohcf':5, 'dohcv':6}
Enco_cp['enginetype'] = Enco_cp['enginetype'].map(name_mapping5)

In [23]: # custom mapping for fuelsystem
name_mapping6 = {'mpfi':1, '2bbl':2, 'mfi':3, '1bbl':4, 'spfi':5, '4bbl':6, 'idi':7, 'sp'}
Enco_cp['fuelsystem'] = Enco_cp['fuelsystem'].map(name_mapping6)

In [24]: Enco_cp.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   Car_Name              205 non-null   int32
3   fueltype              205 non-null   int64
4   aspiration             205 non-null   int64
5   doornumber            205 non-null   int32
6   carbody               205 non-null   int64
7   drivewheel            205 non-null   int64
8   enginelocation         205 non-null   int64
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   int64
15  cylindernumber        205 non-null   int32
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   int64
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm               205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg            205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int32(3), int64(15)
memory usage: 39.4 KB

In [25]: ## syntex for object type columns conversion confiremation
object_cols = [col for col in cp.columns if cp[col].dtype == 'object']
object_cols1 = [col for col in Enco_cp.columns if Enco_cp[col].dtype == 'object']

```

```

if object_cols + object_cols1 == object_cols:
    print('All Object type columns are converted to numerical columns')
else:
    print('Object type columns conversion are pending')

```

All Object type columns are converted to numerical columns

```

In [26]: ## finding the Correlation of Data Frame
cor_Enco_cp = Enco_cp.corr()

```

```


In [27]: ## scaling dataframe
Scal_cp = Enco_cp
scaler = MinMaxScaler()
scaled_data = scaler.fit_transform(Scal_cp)
cp_scaled = pd.DataFrame(scaled_data, columns=Scal_cp.columns)

```

```

In [28]: ## finding correlation of the scaled data frame
cor_scl_cp = cp_scaled.corr()
## Plotting heatmap of Correlation
plt.figure(figsize = (10,8))
sns.heatmap(cor_scl_cp,annot=True,cmap='coolwarm',vmin=-1,vmax=1,annot_kws={"size": 5})
plt.show()

```



```

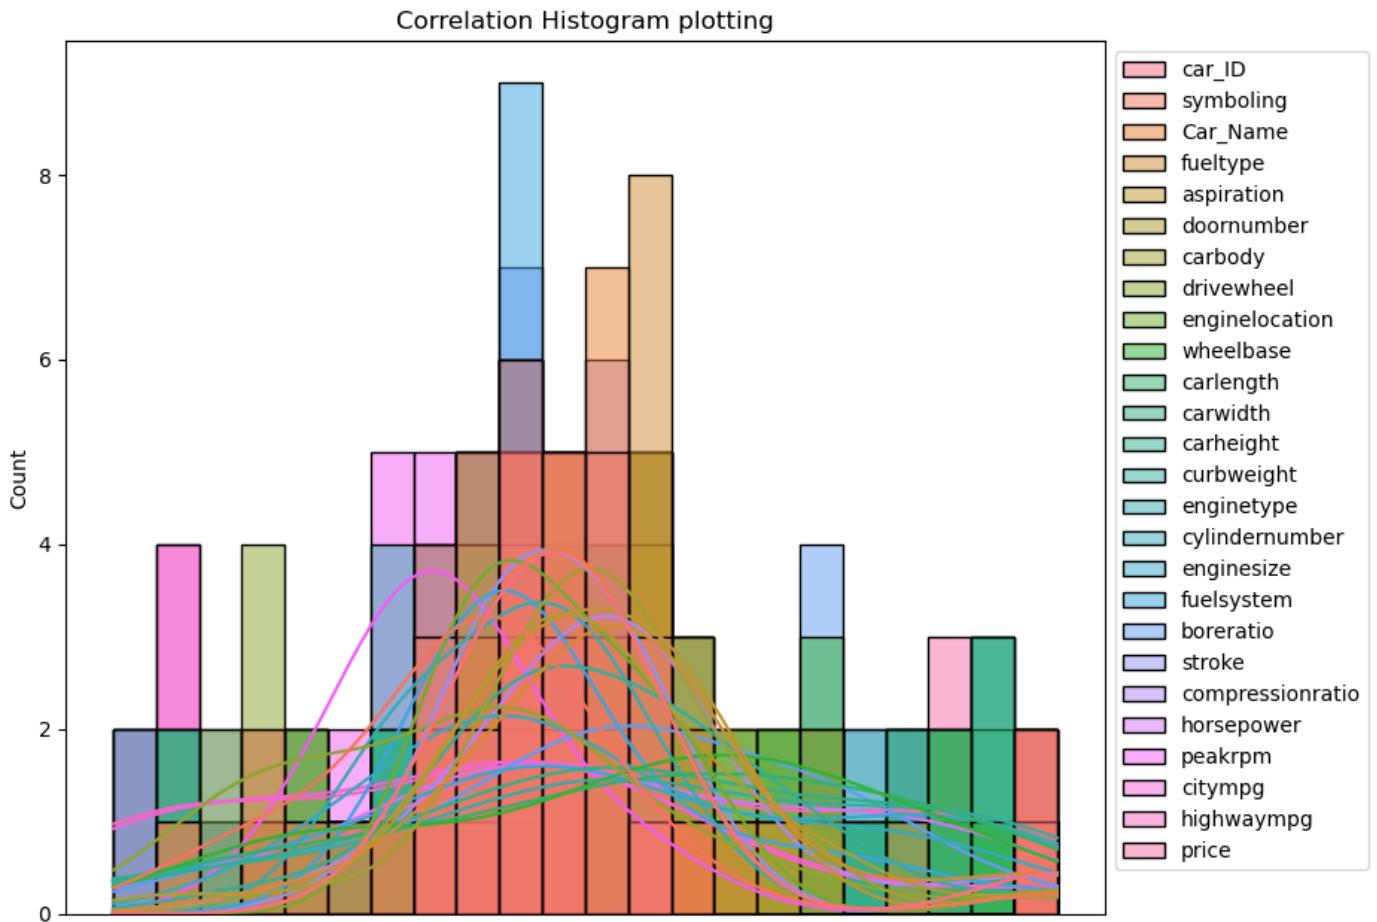
In [29]: # Correlation plotting
plt.figure(figsize=(9, 7.6))
plt.title("Correlation Histogram plotting")

```

```

ax=sns.histplot(data=cor_scl_cp,kde=True)
sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
plt.xticks([])
plt.show()

```



```

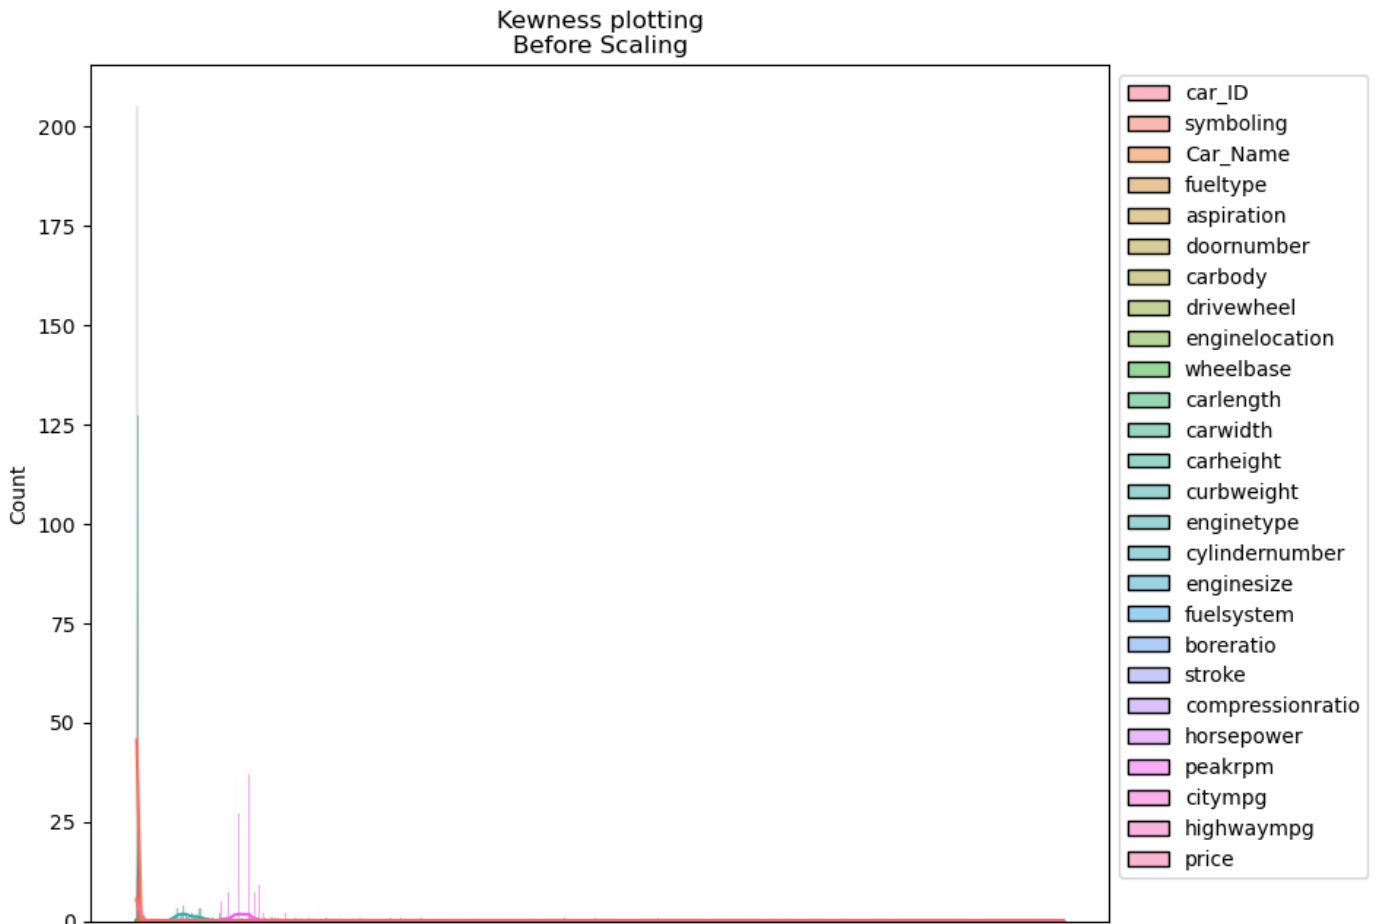
In [30]: skewness = cp_scaled.skew()
print(skewness)


```

car_ID	0.000000
symboling	0.211072
Car_Name	-0.132597
fueltype	2.732619
aspiration	1.673832
doornumber	-0.247552
carbody	-0.415674
drivewheel	0.058352
enginelocation	8.143531
wheelbase	1.050214
carlength	0.155954
carwidth	0.904003
carheight	0.063123
curbweight	0.681398
enginetype	0.390540
cylindernumber	2.817459
enginesize	1.947655
fuelsystem	1.539368
boreratio	0.020156
stroke	-0.689705
compressionratio	2.610862
horsepower	1.405310
peakrpm	0.075159
citympg	0.663704
highwaympg	0.539997
price	1.777678
dtype:	float64

```
In [31]: #Skewness plotting before mapping and scaling
plt.figure(figsize=(9, 7.6))
plt.title("Kewness plotting\nBefore Scaling")
ax=sns.histplot(data=Enco_cp,kde=True)
sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
plt.xticks([])
plt.show()

#skewness after mapping and scaling
plt.figure(figsize=(9, 7.6))
plt.title("Skewness plotting\nAfter Scaling")
ax=sns.histplot(data=cp_scaled,kde=True)
sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
plt.xticks([])
plt.show()
```





Removing outliers using IQR method

```
In [33]: #IQR Method
def outliers(df):
    for col in df.select_dtypes(include=['int64', 'float64']).columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1

        lower = Q1 - (1.5*IQR)
        upper = Q3 + (1.5*IQR)

        # Capping
        df[col] = df[col].apply(lambda x: lower if x < lower else upper if x > upper else x)

    return df
```

Renaming final data

```
In [35]: cp_scaled.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   car_ID                205 non-null   float64
1   symboling              205 non-null   float64
2   Car_Name              205 non-null   float64
3   fueltype              205 non-null   float64
4   aspiration            205 non-null   float64
5   doornumber            205 non-null   float64
6   carbody               205 non-null   float64
```

```

7   drivewheel      205 non-null    float64
8   enginelocation  205 non-null    float64
9   wheelbase       205 non-null    float64
10  carlength       205 non-null    float64
11  carwidth        205 non-null    float64
12  carheight       205 non-null    float64
13  curbweight      205 non-null    float64
14  enginetype      205 non-null    float64
15  cylindernumber  205 non-null    float64
16  enginesize      205 non-null    float64
17  fuelsystem      205 non-null    float64
18  boreratio       205 non-null    float64
19  stroke          205 non-null    float64
20  compressionratio 205 non-null    float64
21  horsepower      205 non-null    float64
22  peakrpm         205 non-null    float64
23  citympg         205 non-null    float64
24  highwaympg      205 non-null    float64
25  price           205 non-null    float64

```

dtypes: float64(26)

memory usage: 41.8 KB

```
In [36]: Car_price = cp_scaled
```

```
In [37]: Car_price = outliers(Car_price)
Car_price.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 205 entries, 0 to 204



Data columns (total 26 columns):



#	Column	Non-Null Count	Dtype
0	car_ID	205 non-null	float64
1	symboling	205 non-null	float64
2	Car_Name	205 non-null	float64
3	fueltype	205 non-null	float64
4	aspiration	205 non-null	float64
5	doornumber	205 non-null	float64
6	carbody	205 non-null	float64
7	drivewheel	205 non-null	float64
8	enginelocation	205 non-null	float64
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	float64
14	enginetype	205 non-null	float64
15	cylindernumber	205 non-null	float64
16	enginesize	205 non-null	float64
17	fuelsystem	205 non-null	float64
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	float64
22	peakrpm	205 non-null	float64
23	citympg	205 non-null	float64
24	highwaympg	205 non-null	float64
25	price	205 non-null	float64



dtypes: float64(26)



memory usage: 41.8 KB



```
In [38]: ## plotting boxplot to visualise removed outliers
for i in Car_price.columns:
    sns.boxplot(data=Car_price,x=i)
plt.show()
```














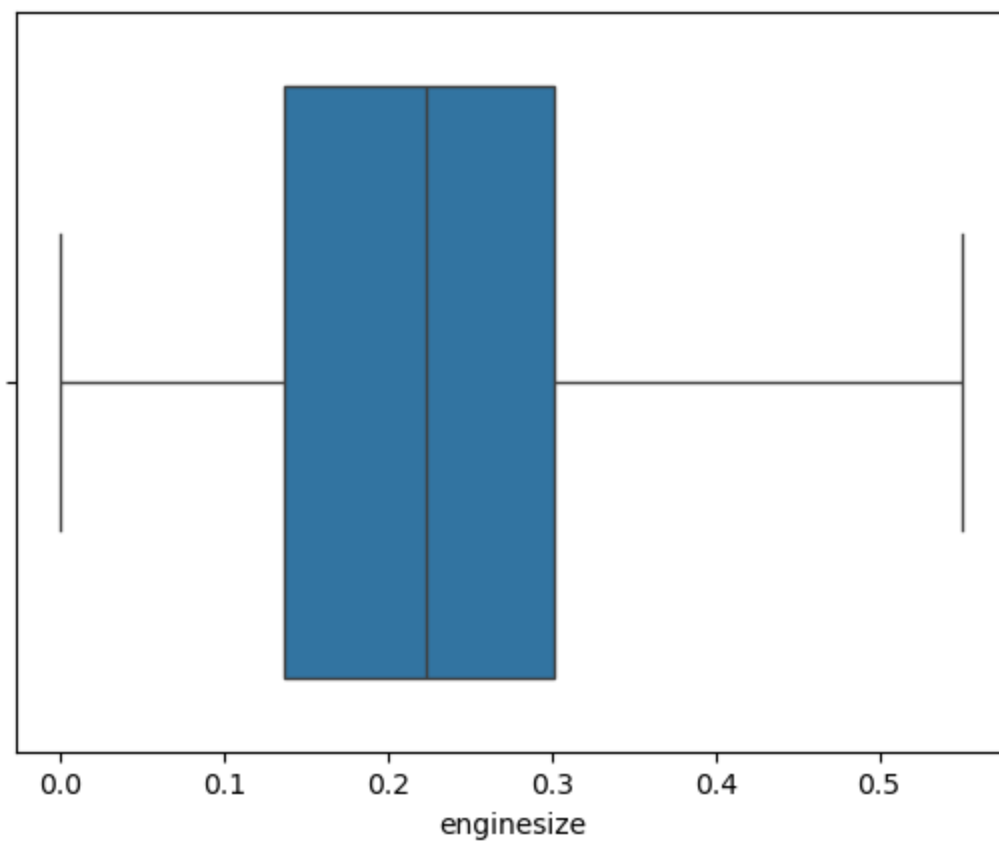




























2.Model Implementation

In []:

```
In [40]: # Assainginn x value and y value as Fetures and target  
x=cp_scaled.drop(columns='price')  
y=cp_scaled['price']
```

1)Linear Regression

```
In [42]: # syntax of training and testing for linear regression
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=5)
```

```
In [43]: from sklearn.linear_model import LinearRegression
#syntax for linear regression modeling
model = LinearRegression()
model.fit(x_train, y_train)
```

```
Out[43]: ▼ LinearRegression ⓘ ?
LinearRegression()
```

```
In [44]: # Syntax for model Prediction
y_pred = model.predict(x_test)
print(y_pred)
```

```
[0.0173999  0.04114517 0.26544993 0.51049577 0.08589462 0.24499946
 0.36402258 0.07779068 0.04451748 0.03283352 0.0672954  0.07635405
 0.07446685 0.2553732  0.09497911 0.08987182 0.22474338 0.53989787
 0.03839838 0.02304198 0.01413171 0.13812371 0.13675742 0.33336011
 0.11148947 0.17399271 0.2307172  0.04856993 0.29493778 0.052302
 0.08707402 0.01230129 0.30648486 0.1690059  0.11836766 0.12674929
 0.14777943 0.60542014 0.46073699 0.05247625 0.07869153]
```

```
In [45]: # Syntax for model score
score_LinearRegression = model.score(x_test,y_test)
score_LinearRegression
```

```
Out[45]: 0.9218232858292736
```

2)Decision Tree Regressor

```
In [47]: #syntax of training and testing for Decision Tree Regressor
x1_train, x1_test, y1_train, y1_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [48]: #syntax for Decision Tree Regressor modeling
from sklearn.tree import DecisionTreeRegressor
modell = DecisionTreeRegressor(random_state=42)
modell.fit(x1_train, y1_train)
```

```
Out[48]: ▼ DecisionTreeRegressor ⓘ ?
DecisionTreeRegressor(random_state=42)
```

```
In [49]: #syntax for model prediction
y1_pred = modell.predict(x1_test)
print(y1_pred)
```

```
[0.60715704 0.28399782 0.09510451 0.17506579 0.60715704 0.00191152
 0.04468497 0.08043295 0.09440941 0.07142148 0.17506579 0.07298545
 0.21168264 0.15217715 0.60715704 0.03028648 0.00697582 0.18981183
 0.04644754 0.09510451 0.12727769 0.23040068 0.05910829 0.00672757
 0.06183903 0.60715704 0.08269202 0.28293034 0.05910829 0.2697731
 0.60715704 0.03850355 0.04138325 0.34573755 0.07047813 0.60715704
 0.16324909 0.16699767 0.0942977  0.24206842 0.10865895]
```

```
In [50]: #syntax for model score
```



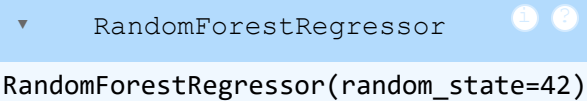
```
score_DecisionTreeRegressor = model1.score(x1_test,y1_test)
score_DecisionTreeRegressor
```

Out[50]: 0.9545150840364851

3) RandomForestRegressor

In [52]: *#syntax of traning and testing for RandomForestRegressor*
`x2_train,x2_test,y2_train,y2_test=train_test_split(x, y, test_size=0.2,random_state=42)`

In [53]: *#syntas for RandomForestRegressor modeling*
`from sklearn.ensemble import RandomForestRegressor`
`model2=RandomForestRegressor(random_state=42)`
`model2.fit(x2_train,y2_train)`

Out[53]: 
RandomForestRegressor(random_state=42)

In [54]: *# syntax of model prediction*
`y2_pred = model2.predict(x2_test)`
`print(y2_pred)`

```
[0.60248101 0.34357182 0.09436721 0.19631374 0.50423291 0.03565563
 0.06862867 0.07303436 0.10732858 0.07418524 0.21084082 0.07193039
 0.21741274 0.1443836 0.60715704 0.03391614 0.01339506 0.22582171
 0.08858175 0.09936547 0.13185939 0.24552319 0.05238792 0.01416836
 0.05285115 0.60574189 0.09671913 0.28936448 0.05539943 0.28403158
 0.50567524 0.03384365 0.07140212 0.348156 0.07165384 0.50048856
 0.11344124 0.18313688 0.0584498 0.23754655 0.07887121]
```

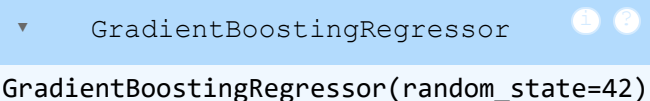
In [55]: *# syntax of model score*
`score_RandomForestRegressor = model2.score(x2_test,y2_test)`
`print(score_RandomForestRegressor)`

0.9568274904378331

4) GradientBoostingRegressor

In [57]: *# Syntax of traning and testing for Gradient boosting regressor*
`x3_train,x3_test,y3_train,y3_test = train_test_split(x,y,test_size=0.2,random_state=42)`

In [58]: *# syntax for GradientboostingRegressor*
`from sklearn.ensemble import GradientBoostingRegressor`
`model3 = GradientBoostingRegressor(random_state = 42)`
`model3.fit(x3_train,y3_train)`

Out[58]: 
GradientBoostingRegressor(random_state=42)

In [59]: *# Syntax of model prediction*
`y3_pred = model.predict(x3_test)`
`print(y3_pred)`

```
[ 0.50408979 0.31074129 0.10535158 0.199467 0.51442429 0.01230129
 0.04451748 0.03366744 0.13834713 0.08050286 0.1690059 0.05779177
 0.28184428 0.10256637 0.55473401 0.01413171 -0.06475342 0.21771767
 0.10351905 0.13108738 0.13971486 0.36231498 0.04322966 0.01511448
 0.07779068 0.50523767 0.19733959 0.27348636 0.02185689 0.24499946]
```

```
0.47402081 0.01740523 0.0173999 0.37884679 0.04590541 0.47502946
0.12009852 0.14921605 0.06041256 0.22474338 0.09772654]
```

```
In [60]: # syntax for model score
score_GradientBoostingRegressor = model3.score(x3_test,y3_test)
print(score_GradientBoostingRegressor)

0.9654483345578171
```

Support Vector Regressor(SVR)

```
In [62]: # syntax of standardscaling of x and y value
scaler_x = StandardScaler()
x_scaled = scaler_x.fit_transform(x)
scaler_y = StandardScaler()
y_scaled = scaler_y.fit_transform(y.values.reshape(-1,1))
```

```
In [63]: # syntax of traing and testng for SVR
x4_train,x4_test,y4_train,y4_test = train_test_split(x_scaled,y_scaled, test_size=0.2,ra
```

```
In [64]: # Syntax of DVR modeling
from sklearn.svm import SVR
model4 = SVR(kernel='rbf',C=1.0, epsilon=0.1)
model4.fit(x4_train,y4_train.ravel())
```

```
Out[64]:
```

▼ SVR ⓘ ?

SVR()

```
In [65]: # sytax of model predictin
y4_pred = model4.predict(x4_test)
print(y4_pred)

[ 1.70411533  0.59824683 -0.86953452  0.21991343  1.41229274 -1.08769008
 -0.44380025 -0.93592188 -0.45225855 -0.59721948 -0.08104726 -0.76734628
  0.38416197 -0.14378696  2.49010599 -0.86407123 -0.22082264  0.13084676
 -0.77148585 -0.4974642  -0.51203069  0.42334069 -0.54468937 -0.46899712
 -0.81044879  1.8999882  -0.36339444  0.4711522  -0.94643054  0.50326225
  1.56182218 -0.86067676 -0.62876583  1.01176453 -0.75934718  1.39070537
 -0.3209201  0.02335724 -0.89973135  0.15776104 -0.76007862]
```

```
In [66]: #sytex for model score
score_SVR = model4.score(x4_test,y4_test)
print(score_SVR)

0.8812822630479982
```

```
In [67]: # List of values
values = {'LinearRegression':[score_LinearRegression],
          'DecisionTreeRegressor':[score_DecisionTreeRegressor],
          'RandomForestRegressor':[score_RandomForestRegressor],
          'GradientBoostingRegressor':[score_GradientBoostingRegressor],
          'SVR':[score_SVR]}
df_values=pd.DataFrame(values)
max_values = df_values.max()
overall_max_value = df_values.max().max()
overall_max_value
overall_max_columns = df_values.columns[df_values.eq(overall_max_value).any()]
print(f'The highest model score of Regression model is {overall_max_columns}\nAnd the Sc

The highest model score of Regression model is Index(['GradientBoostingRegressor'], dtype='object')
Andthe Score is 0.9654483345578171
```

3). Model Evelution

Comparing the performance of all models based on R-squared, Mean Squared Error (MSE) and Mean Absolute Error (MAE)

```
In [69]: from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
```

Evaluating Linear Regression model using following metrics

```
In [71]: mae_LR = mean_absolute_error(y_test, y_pred)
mse_LR = mean_squared_error(y_test, y_pred)
r2_LR = r2_score(y_test, y_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Linear Regres
print_section(f'Mean Squared Error: {mse_LR}\nMean Absolute Error: {mae_LR}\nR-Squared S
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Linear Regression model
-----
```

```
Mean Squared Error: 0.0023150850976519604
Mean Absolute Error: 0.03594446242220855
R-Squared Score: 0.9218232858292736
-----
```

Evaluating Decision Tree Regressor model using following metrics

```
In [73]: mae_DTR = mean_absolute_error(y1_test, y1_pred)
mse_DTR = mean_squared_error(y1_test, y1_pred)
r2_DTR = r2_score(y1_test, y1_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Decidion Tree
print_section(f'Mean Squared Error: {mse_DTR}\nMean Absolute Error: {mae_DTR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Decidion Tree Regresso model
-----
```

```
Mean Squared Error: 0.0014845947267692608
Mean Absolute Error: 0.027874622327227223
R-Squared Score: 0.9545150840364851
-----
```

Evaluating Random Forest Regressor model using following metrics

```
In [75]: mae_RFR = mean_absolute_error(y2_test, y2_pred)
mse_RFR = mean_squared_error(y2_test, y2_pred)
r2_RFR = r2_score(y2_test, y2_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Random Forest
print_section(f'Mean Squared Error: {mse_RFR}\nMean Absolute Error: {mae_RFR}\nR-Squared
```

```
-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Random Forest Regressor model
-----
```

```
Mean Squared Error: 0.00140911945597086
Mean Absolute Error: 0.026968862204385824
R-Squared Score: 0.9568274904378331
-----
```

Evaluating Gradient Boosting Regressor model using following metrics

```
In [77]: mae_GBR = mean_absolute_error(y3_test, y3_pred)
mse_GBR = mean_squared_error(y3_test, y3_pred)
```

```

r2_GBR = r2_score(y3_test,y3_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Gradient Boos
print_section(f'Mean Squared Error: {mse_GBR}\nMean Absolute Error: {mae_GBR}\nR-Squared

```

```

-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Gradient Boosting Regressor model
-----

```

```

Mean Squared Error: 0.002743681669317043
Mean Absolute Error: 0.0418180493930531
R-Squared Score: 0.9159392607899799
-----

```

Evaluating Support Vector Regressor model using following metrics

```

In [79]: mae_SVR = mean_absolute_error(y4_test,y4_pred)
mse_SVR = mean_squared_error(y4_test,y4_pred)
r2_SVR = r2_score(y4_test,y4_pred)
print_title(f'The Evaluated performace of MSE, MAE and R-Squared Score\nfor Support Vecto
print_section(f'Mean Squared Error: {mse_SVR}\nMean Absolute Error: {mae_SVR}\nR-Squared

```

```

-----
The Evaluated performace of MSE, MAE and R-Squared Score
for Support Vector Regressor model
-----

```

```

Mean Squared Error: 0.14005505830429088
Mean Absolute Error: 0.28534068039624766
R-Squared Score: 0.8812822630479982
-----

```

Best Performing model with justification

```

In [81]: R_score = {'model':['Linear Regression','Decidion Tree Regressor','Random Forest Regress
          'Gradient Boosting Regressor','Support Vector Regressor'],
          'R2_Score %':[100*r2_LR,100*r2_DTR,100*r2_RFR,100*r2_GBR,100*r2_SVR]}
R_score = pd.DataFrame(R_score)
high= R_score['R2_Score %'].max()
low= R_score['R2_Score %'].min()
Iden_model = R_score.loc[R_score['R2_Score %'] == high, 'model'].item()
wrost_model = R_score.loc[R_score['R2_Score %'] == low, 'model'].item()
print_title('          "Result Comparison using R-Squared Score")
print_section(f'Compared and Identifeyed model using R2_Score metric\nof \033[1m{Iden_mod

```

```

-----
Result Comparison using R-Squared Score
-----

```

```

Compared and Identifeyed model using R2_Score metric
of Random Forest Regressor its R-Squared Score
is 95% and it has best-performing algoritham
wtih justification. And the Wroست model
is Support Vector Regressor its R-Squared score is 88%
-----

```

4)Feature Importance Analysis

Identifying the significant variables affecting car prices


```

In [84]: # random forest for feature importance on a regression problem
from sklearn.datasets import make_regression
importance = model2.feature_importances_
# summarize feature importance
for i,v in enumerate(importance):
    print('Feature: %0d, Score: %.5f' % (i,v))

```

```
Feature: 0, Score: 0.01921
Feature: 1, Score: 0.00098
Feature: 2, Score: 0.01325
Feature: 3, Score: 0.00000
Feature: 4, Score: 0.00000
Feature: 5, Score: 0.00055
Feature: 6, Score: 0.00168
Feature: 7, Score: 0.00197
Feature: 8, Score: 0.00000
Feature: 9, Score: 0.00653
Feature: 10, Score: 0.01077
Feature: 11, Score: 0.01080
Feature: 12, Score: 0.00467
Feature: 13, Score: 0.53391
Feature: 14, Score: 0.00000
Feature: 15, Score: 0.00000
Feature: 16, Score: 0.24616
Feature: 17, Score: 0.00286
Feature: 18, Score: 0.00519
Feature: 19, Score: 0.00444
Feature: 20, Score: 0.00250
Feature: 21, Score: 0.04080
Feature: 22, Score: 0.00554
Feature: 23, Score: 0.02291
Feature: 24, Score: 0.06529
```

```
In [85]: variables = list(x.columns.values)
variables
dic = dict(zip(variables,importance))
df_dic = pd.DataFrame([dic])
max_value = max(dic, key=dic.get)
# plot feature importance
plt.bar(variables,importance)
plt.title('Plotting Significat variables')
plt.xlabel('Variables')
plt.xticks(rotation=90, ha='right')
plt.ylabel('Values')
plt.show()
```



5) Hyperparameter Tuning

```
In [228...] from sklearn import svm
from sklearn.svm import SVR
from sklearn.model_selection import cross_val_score
```

```
In [217...] x5_train, x5_test, y5_train, y5_test = train_test_split(x,y,test_size=0.3)
```

```
In [219...] model5 = svm.SVR(kernel='rbf',C=30,gamma='auto')
model5.fit(x5_train, y5_train)
```

```
Out[219]: SVR
SVR(C=30, gamma='auto')
```

```
In [240...] kernels = ['rbf','linear']
C = [1,10,20]
avg_scores = {}
for kval in kernels:
    for cval in C:
        cv_scores = cross_val_score(svm.SVR(kernel=kval,C=cval,gamma='auto'),x,y,cv=5)
        avg_scores[kval + '_' + str(cval)] = np.average(cv_scores)
```

```
In [238...] avg_scores
```

```
Out[238]: {'rbf_1': 0.5987282035895402,  
          'rbf_10': 0.6152560732739006,  
          'rbf_20': 0.6140176596351637,  
          'linear_1': 0.5885914453891583,  
          'linear_10': 0.5690558781453141,  
          'linear_20': 0.5715248726425439}
```